

000
001
002
003
004
005
006
007
008
009
010
011
012
013
014
015
016
017
018
019
020
021
022
023
024
025
026
027
028
029
030
031
032
033
034
035
036
037
038
039
040
041
042
043
044
045
046
047
048
049
050
051
052
053

GEOODE-KD: DISTILLING SENTENCE EMBEDDINGS AS TEACHER-GUIDED GEOMETRIC DYNAMICS

Anonymous authors
Paper under double-blind review

ABSTRACT

Table 1: Main results for distillation from Qwen3-Embedding-4B to BERT-base. Classification datasets are evaluated with F1, pair-classification datasets with average precision (AP), and semantic textual similarity datasets with Spearman correlation. $\star$ denotes an in-domain benchmark. Higher is better.

Method	Classification (F1)			Pair Classification (AP)			STS (Spearman)			Avg.
	Banking77	Tweet	Emotion*	MRPC	SciTail	WiC*	SICK	STS12	STS-B*	
<i>(a) Qwen3-Embedding-4B $\rightarrow$ BERT-base (109M)</i>										
Teacher	93.63	72.22	68.94	84.79	91.77	66.80	83.43	82.55	86.87	81.22
Student base	84.86	47.79	68.02	77.36	67.74	59.22	42.43	21.54	20.30	54.36
Stella	91.08	73.40	64.40	86.23	83.22	68.98	78.38	71.31	79.16	77.35
CDM	89.71	70.76	52.67	84.03	71.47	69.83	71.36	68.87	74.68	72.60
DSKD	88.89	69.31	53.39	86.35	67.73	68.89	70.58	67.99	74.29	71.94
EMO	85.92	69.39	55.41	85.79	64.34	67.73	63.79	64.41	70.14	69.60
TALAS	93.34	74.45	72.70	86.16	87.78	68.11	80.70	77.34	83.57	80.46
GeoODE-KD	92.92	74.05	72.49	85.89	86.12	66.60	80.78	76.01	81.23	79.57

108	1	INTRODUCTION
109		
110	2	DISTILLATION SHOULD SPECIFY A PATH, NOT ONLY AN ENDPOINT
111		
112	2.1	EMBEDDING DISTILLATION AS STATIC STATE MATCHING
113		
114	2.2	REPRESENTATION DEPTH AS SEMANTIC TIME
115		
116	2.3	CONNECTIONS TO LAYER-WISE, GEOMETRIC, AND FLOW-BASED DISTILLATION
117		
118	3	GEOODE-KD: FROM TEACHER EMBEDDINGS TO SEMANTIC DYNAMICS
119		
120	3.1	CONSTRUCTING A GEOMETRY-FAITHFUL TEACHER SPACE
121	3.2	TURNING SEMANTIC DISCREPANCY INTO A VECTOR FIELD
122		
123	3.3	BUILDING A FINITE-HORIZON TRAJECTORY
124		
125	3.4	THE STUDENT TRANSFORMER AS THE FLOW INTEGRATOR
126		
127	3.5	LEARNING OBJECTIVE AND PROPERTIES
128		
129	4	EXPERIMENTS
130		
131	4.1	EXPERIMENTAL SETUP
132	4.2	DOES LEARNING A TRAJECTORY PRODUCE BETTER EMBEDDINGS?
133		
134	4.3	DO STUDENT LAYERS FOLLOW THE PRESCRIBED SEMANTIC FLOW?
135		
136	4.4	WHAT MAKES GEOODE-KD WORK?
137		
138	5	DISCUSSION AND LIMITATIONS
139		
140	6	CONCLUSION
141		
142		AI USE STATEMENT
143		ETHICS STATEMENT
144		REPRODUCIBILITY STATEMENT
145		
146		REFERENCES
147		
148	A	MATHEMATICAL DERIVATIONS AND PROOFS
149		
150	A.1	RIEMANNIAN GEOMETRY AND VECTOR FIELD DERIVATIONS
151		
152	A.2	FINITE-HORIZON FLOW DERIVATION
153		
154	A.3	PROOFS OF THE FLOW PROPERTIES
155		
156	B	TARGET CONSTRUCTION AND GAUGE ALIGNMENT
157		
158	B.1	PCA INITIALIZATION
159		
160	B.2	GRAM-CONSISTENT TARGET FITTING
161		
	B.3	ORTHOGONAL GAUGE ALIGNMENT
	C	TRAINING ALGORITHM AND IMPLEMENTATION DETAILS
	D	DATASETS, BASELINES, AND EVALUATION PROTOCOL